

Observing Cloud Resources

SRE Project Template <Eslam Adel - eslam.adel.me@gmail.com>

Categorize Responsibilities

Prometheus and Grafana Screenshots


Provide a screenshot of the Prometheus node_exporter service running on the EC2 instance. Use the following command to show that the system is running: `sudo systemctl status node_exporter`

```
ubuntu@ip-172-31-17-52: ~  
ubuntu@ip-172-31-17-52:~$ sudo systemctl status node_exporter  
● node_exporter.service - Node Exporter  
   Loaded: loaded (/etc/systemd/system/node_exporter.service; bad; vendor preset: enabled)  
   Active: active (running) since Thu 2022-12-15 17:31:11 UTC; 21s ago  
     Main PID: 3524 (node_exporter)  
       Tasks: 4 (limit: 1109)  
      CGroup: /system.slice/node_exporter.service  
              └─3524 /usr/local/bin/node_exporter  
  
Dec 15 17:31:11 ip-172-31-17-52 node_exporter[3524]: level=info ts=2022-12-15T17:31:11.140Z caller=node_exporter.go:115 collector=time  
Dec 15 17:31:11 ip-172-31-17-52 node_exporter[3524]: level=info ts=2022-12-15T17:31:11.140Z caller=node_exporter.go:115 collector=timex  
Dec 15 17:31:11 ip-172-31-17-52 node_exporter[3524]: level=info ts=2022-12-15T17:31:11.140Z caller=node_exporter.go:115 collector=udp_queues  
Dec 15 17:31:11 ip-172-31-17-52 node_exporter[3524]: level=info ts=2022-12-15T17:31:11.140Z caller=node_exporter.go:115 collector=uname  
Dec 15 17:31:11 ip-172-31-17-52 node_exporter[3524]: level=info ts=2022-12-15T17:31:11.140Z caller=node_exporter.go:115 collector=vmstat  
Dec 15 17:31:11 ip-172-31-17-52 node_exporter[3524]: level=info ts=2022-12-15T17:31:11.140Z caller=node_exporter.go:115 collector=xfs  
Dec 15 17:31:11 ip-172-31-17-52 node_exporter[3524]: level=info ts=2022-12-15T17:31:11.140Z caller=node_exporter.go:115 collector=zfs  
Dec 15 17:31:11 ip-172-31-17-52 node_exporter[3524]: level=info ts=2022-12-15T17:31:11.140Z caller=node_exporter.go:199 msg="Listening on  
Dec 15 17:31:33 ip-172-31-17-52 systemd[1]: /etc/systemd/system/node_exporter.service:1: Missing '='.  
ubuntu@ip-172-31-17-52:~$ sudo ufw allow 9100/tcp  
Rules updated  
Rules updated (v6)  
ubuntu@ip-172-31-17-52:~$
```

Host Metric (CPU, RAM, Disk, Network)

Dashboard

Disk





Responsibilities

1. The development team wants to release an emergency hotfix to production. Identify two roles of the SRE team who would be involved in this and why.

Monitoring Engineer: is responsible for creating Dashboards, creating alerts, he is the first one to know if an issue also he manages monitoring rules for infrastructure.

Release Engineer: is responsible to check code and its dependencies, using CI/CD to ensure code could be released or not and execute rollback procedures if necessary.

2. The development team is in the early stages of planning to build a new product. Identify two roles of the SRE team that should be invited to the meeting and why.

Team leader: is responsible to contribute to architecture meetings, create workflows of the team and keep them focused.

The architect: is responsible for documentation, making recommendations for new technologies, create a highly scalable infrastructure.

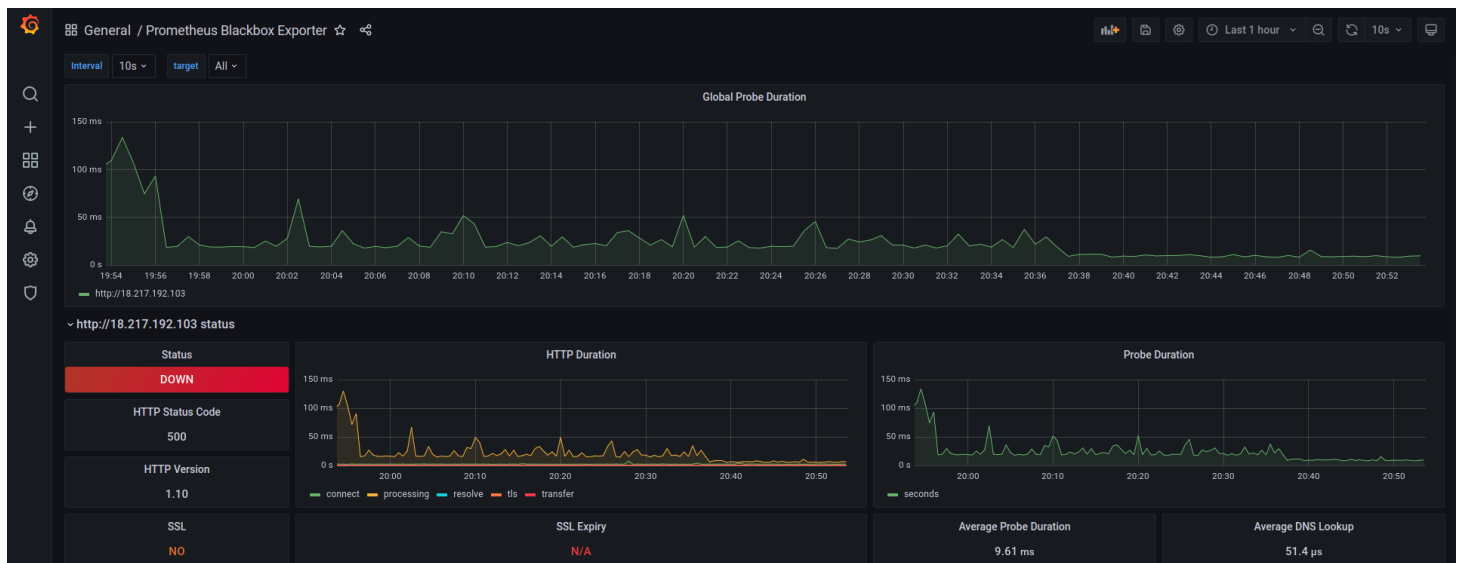
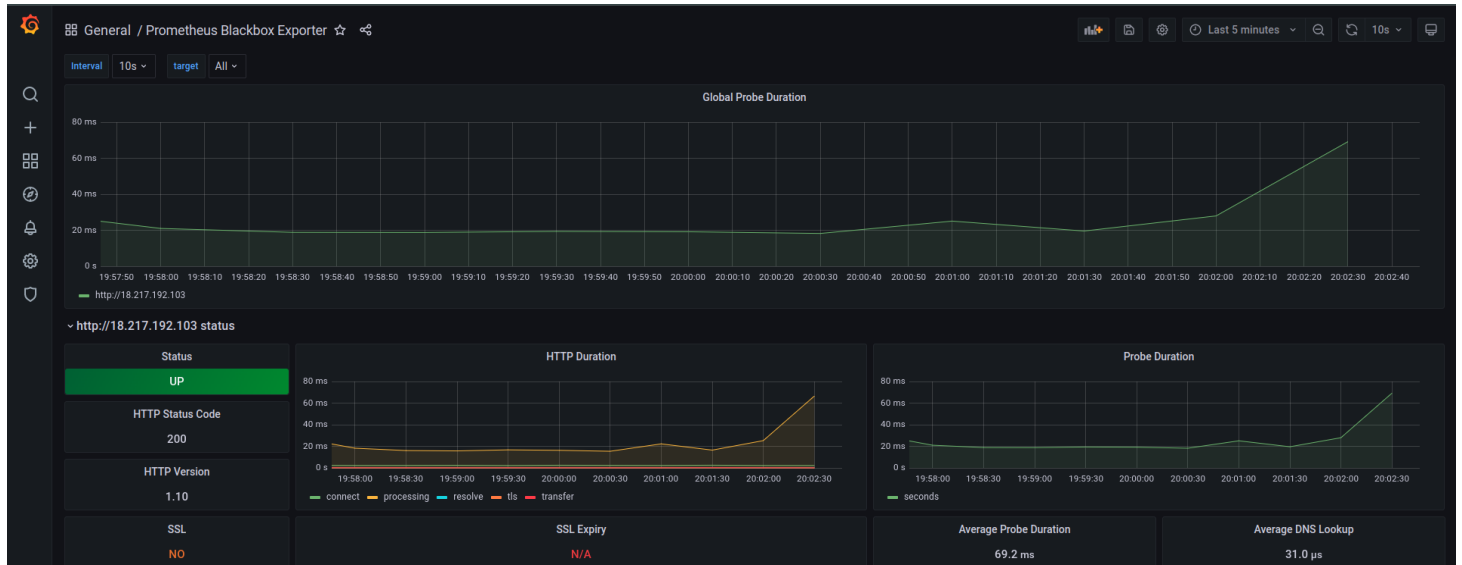
3. The emergency hotfix from question 1 was applied and is causing major issues in production. Which SRE role would primarily be involved in mitigating these issues?


The Release Manager is responsible for handling these situations by running rolling back procedures.

Team Formation and Workflow Identification

API Monitoring and Notifications

Display the status of an API endpoint: Provide a screenshot of the Grafana dashboard that will show at which point the API is unhealthy (non-200 HTTP code), and when it becomes healthy again (200 HTTP code).






Create a notification channel: Provide a screenshot of the Grafana notification which shows the summary of the issue and when it occurred.

High bandwidth above 710
instance:node_network_transmit_bytes:rate:sum{instance="18.217.192.103:9100"}
1274.4582566258

 Grafana v8.1.2 | Today at 8:39 PM


[Alerting] CPU alert
high load CPU above 0.00120
instance:node_cpu:ratio{instance="18.217.192.103:9100", job="ec2"}
0.0041549204810057

 Grafana v8.1.2 | Today at 8:43 PM



incoming-webhook APP 8:48 PM


[Alerting] API health check alert
API is down
probe_http_status_code{instance="http://18.217.192.103", job="blackbox"}
500

 Grafana v8.1.2 | Today at 8:48 PM


[OK] Network alert

 Grafana v8.1.2 | Today at 8:49 PM

[OK] CPU alert

 Grafana v8.1.2 | Today at 8:51 PM

[OK] API health check alert

 Grafana v8.1.2 | Today at 8:55 PM


[Alerting] Network Bytes received alert
instance:node_network_receive_bytes:rate:sum{instance="10.100.3.108:9100"}
27008.921212121

Configure alert rules: Provide a screenshot of the alert rules list in Grafana.

Q Search alerts


StatesAll

How to add an alert

 API health check alert
OK for 2 minutes


Pause

Edit alert

 CPU alert
PENDING for a few seconds


Pause

Edit alert

 Memory alert
OK for 26 minutes


Pause

Edit alert

 Network Bytes received alert
ALERTING for a minute

Pause

Edit alert


 Network alert
OK for 9 minutes

Pause

Edit alert

Applying the Concepts

Graph 1



4a. Given the above graph, where does it show that the API endpoint is down? Where on the graph does this show that the API is healthy again?

At 15:27 API was down, status code was changed from 200 to 500, and it becomes healthy again at 15:37.


4b. If there was no SRE team, how would this outage affect customers?

Will not be able to access the app and the customer would complain about it.

4c. What could be put in place so that the SRE team could know of the outage before the customer does?

By Setting Alerts to send notifications about the problem when they occur to our channels like slack at this project.

Graph 2



5a. Given the above graph, which instance had the increase in traffic, and approximately how many bytes did it receive (feel free to round)?

10.0.0.68:9100 instance, received 3.6k bytes.

5b. Which team members on the SRE team would be interested in this graph and why?

Monitoring Engineers, because its his responsibilities and also he is responsible with:

- *Creating dashboards*
- *Creating and manage alerting rules*
- *He should be the first who knows the issue*
- *He probably use tools like Prometheus/Grafana and ELK.*